
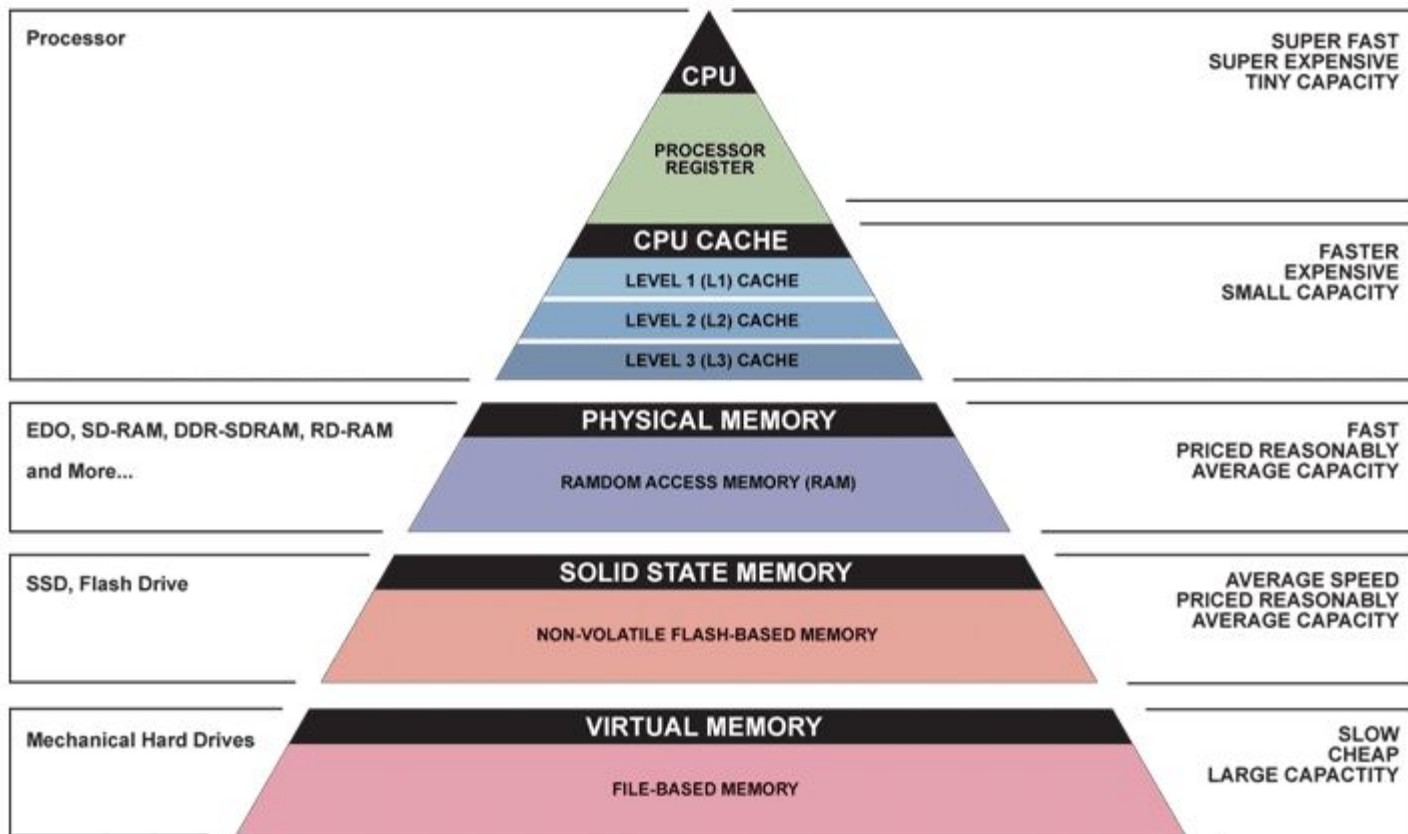


CASS: Exercise session 7

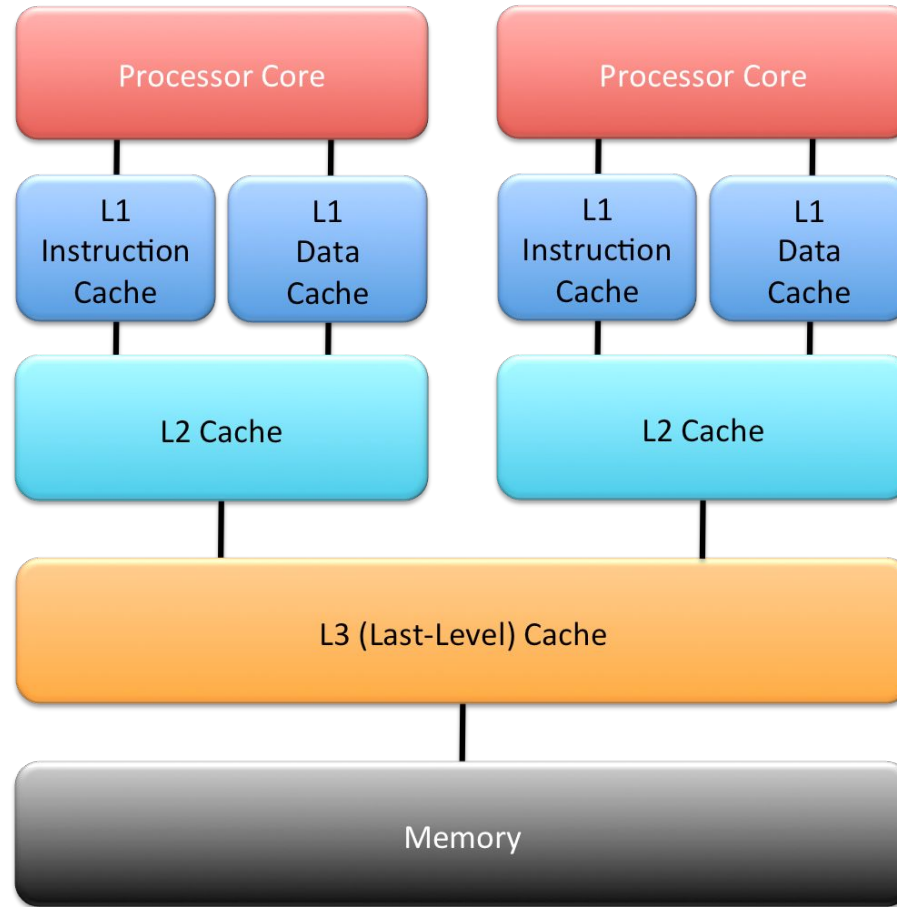
Caches and Microarchitectural Timing Attacks







▲ Simplified Computer Memory Hierarchy
Illustration: Ryan J. Leng



Caches are small.


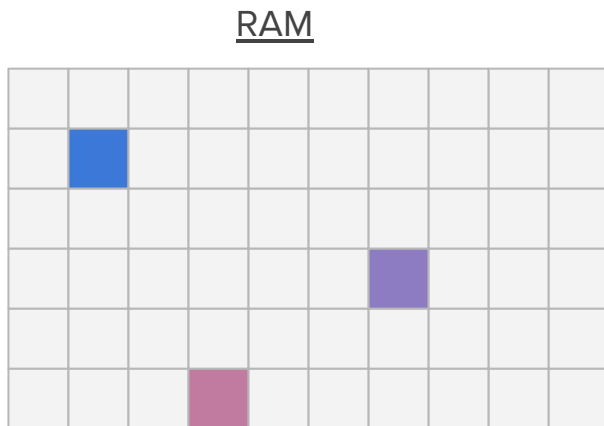
How can they make a difference?

“Locality principle”

Programs access a relatively **small portion** of
address space at a time


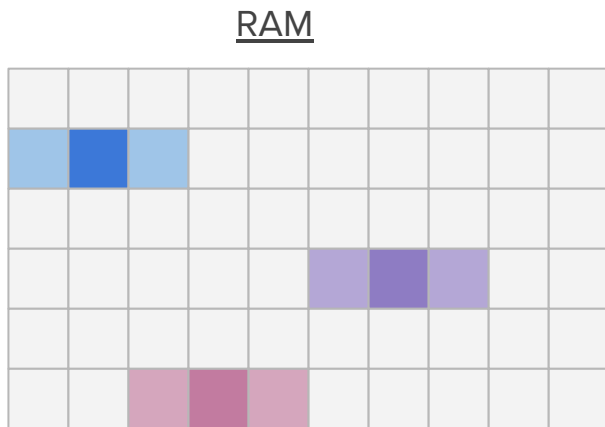
Temporal locality

- Memory location accessed?
 - Same location *likely* accessed again soon!
 - Add it to cache!



Spatial locality

- Memory location accessed?
 - Nearby locations *likely* accessed soon!
- Transfer **block** to cache
 - Cache line = block = #bytes transferred at once

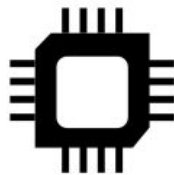


CPU cache operation



Cache principle: CPU speed \gg DRAM latency \rightarrow *cache code/data*

```
while true do  
  maccess(&a);  
endwh
```



CPU + cache






DRAM memory

CPU cache operation



Cache miss: Request data from (slow) DRAM upon first use



CPU + cache

cache miss



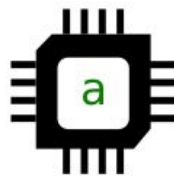
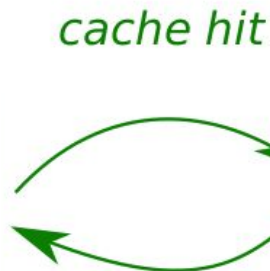
DRAM memory

CPU cache operation



Cache hit: No DRAM access required for subsequent uses

```
while true do  
  maccess(&a);  
endwh
```



CPU + cache



DRAM memory

Cache performance

$$\text{miss rate} = (1 - \text{hit rate})$$

Conclusion:

Small amount of memory, used in smart way






Can we abuse timing differences as an attacker to infer secret information? :-)

CPU's offer **instructions to**
flush the cache!

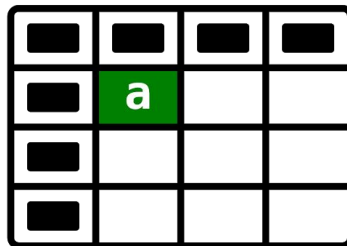


Flush+Reload: Cache timing attacks on shared memory



```
if secret do
    maccess(&a);
else
    maccess(&b);
endif
```


*'a' is accessible
to attacker*



CPU cache

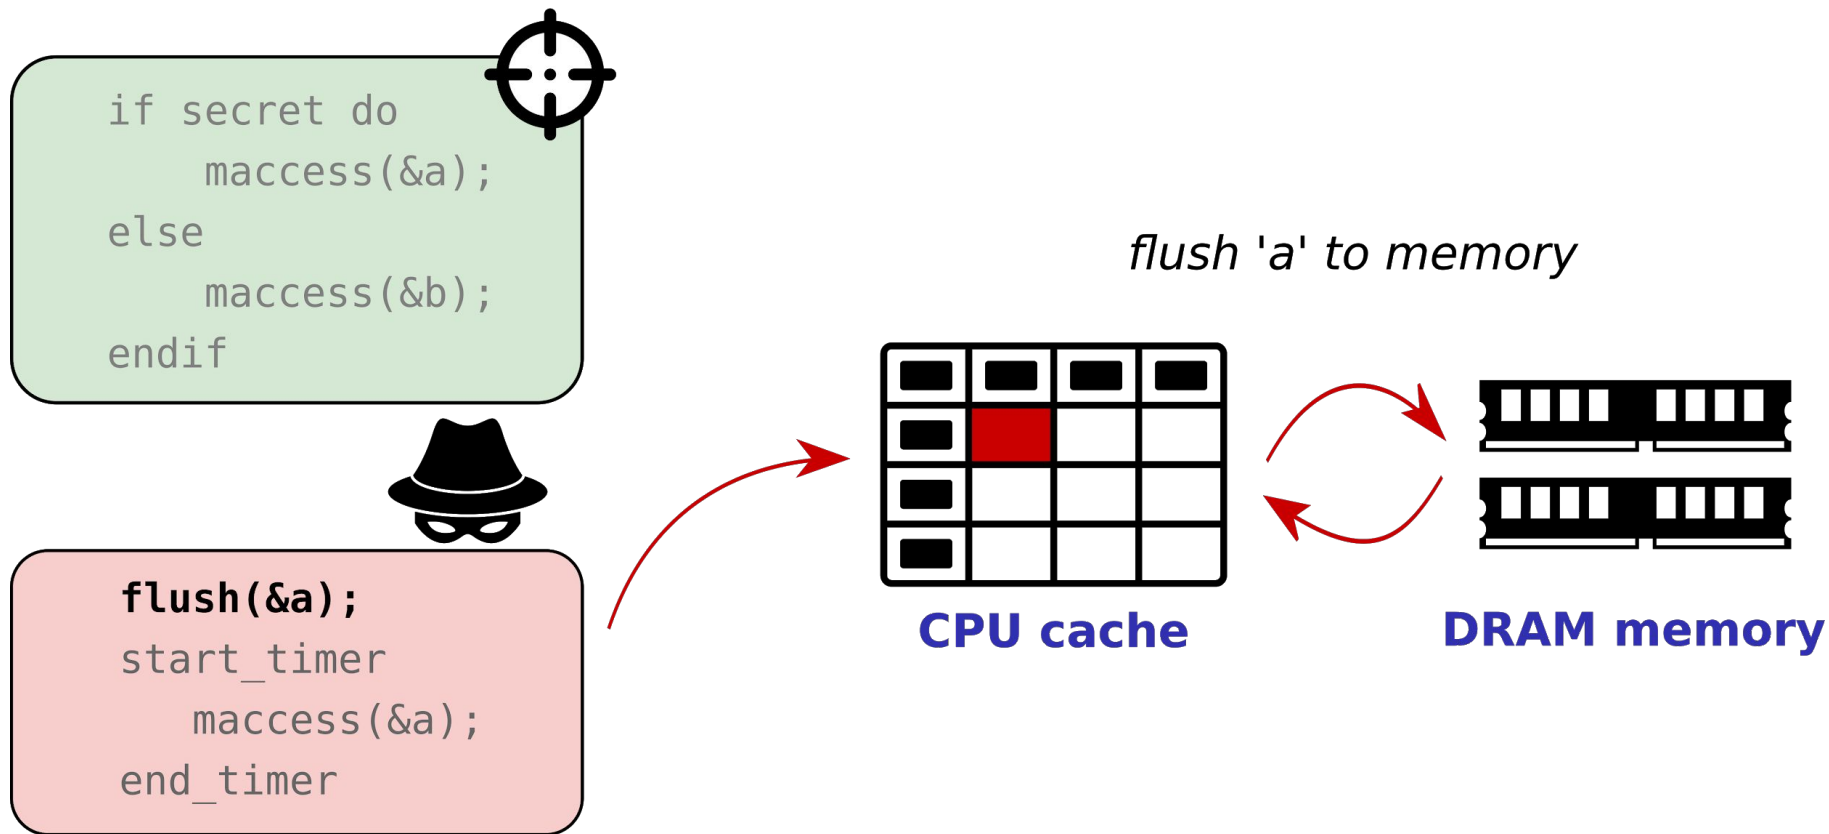


DRAM memory




```
flush(&a);
start_timer
maccess(&a);
end_timer
```

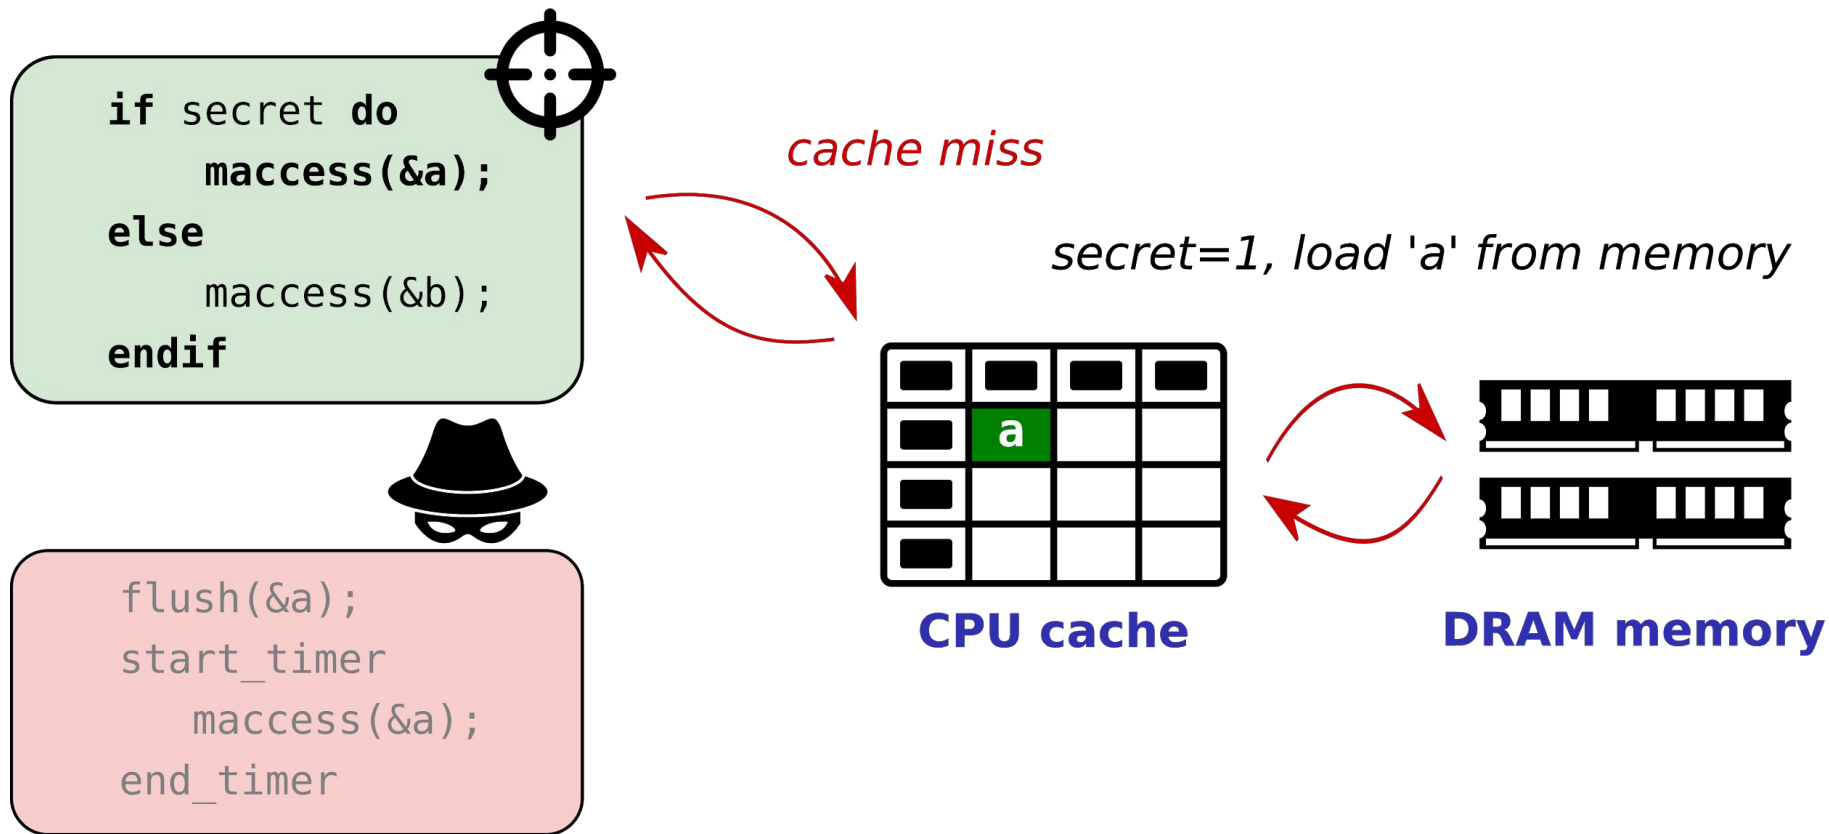
Flush+Reload: Cache timing attacks on shared memory



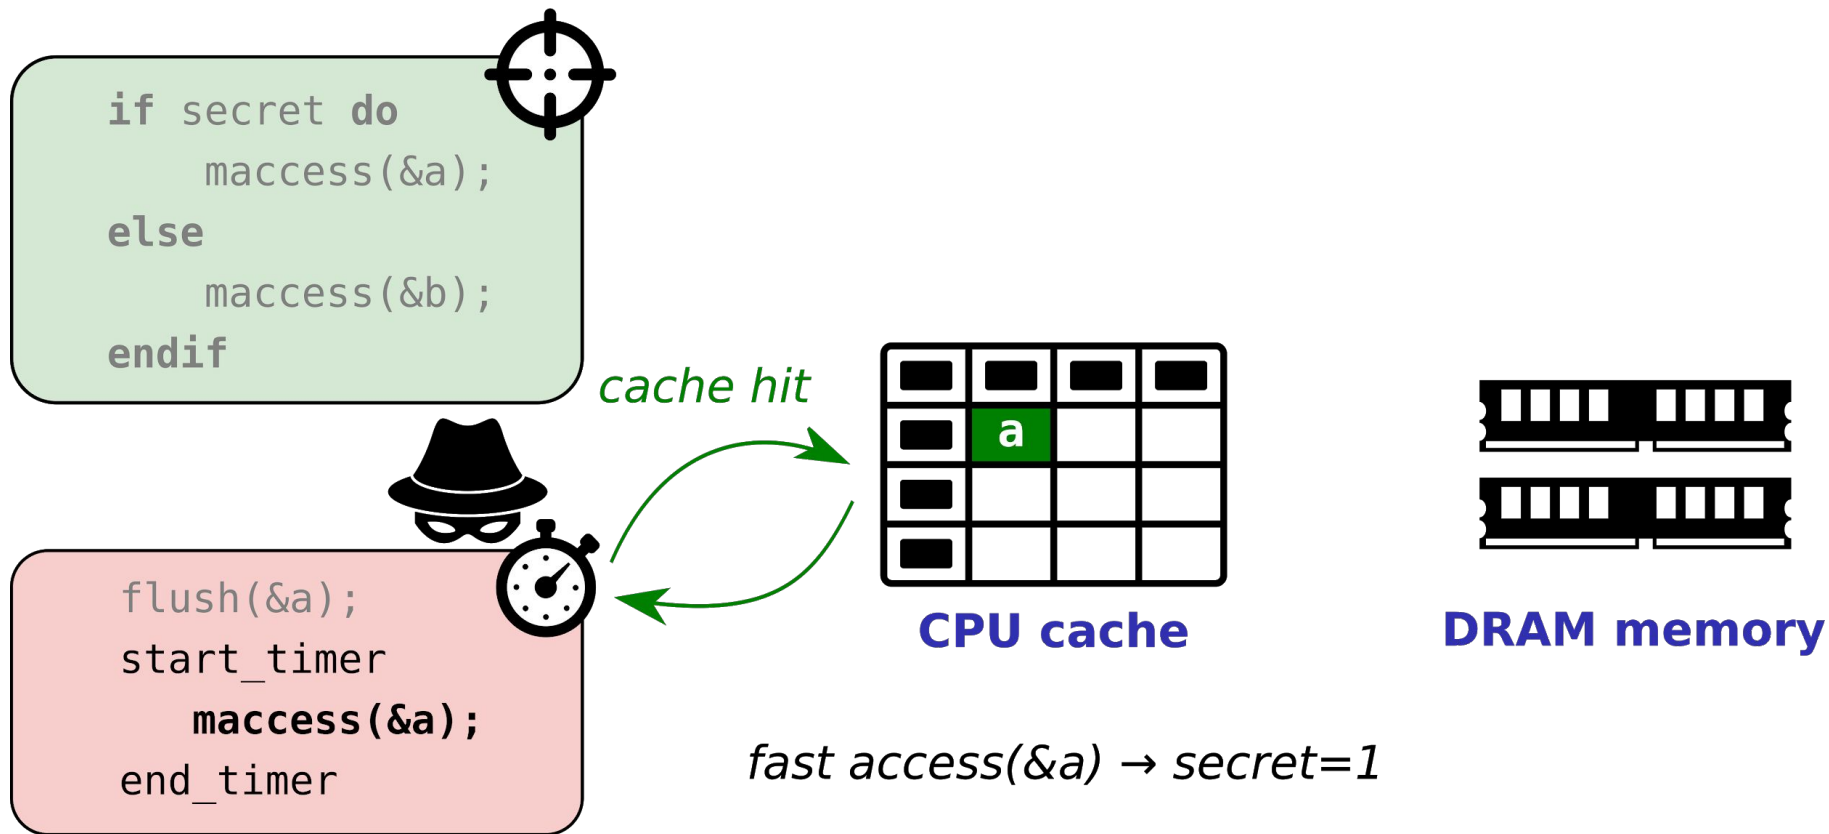
Flush+Reload: Cache timing attacks on shared memory



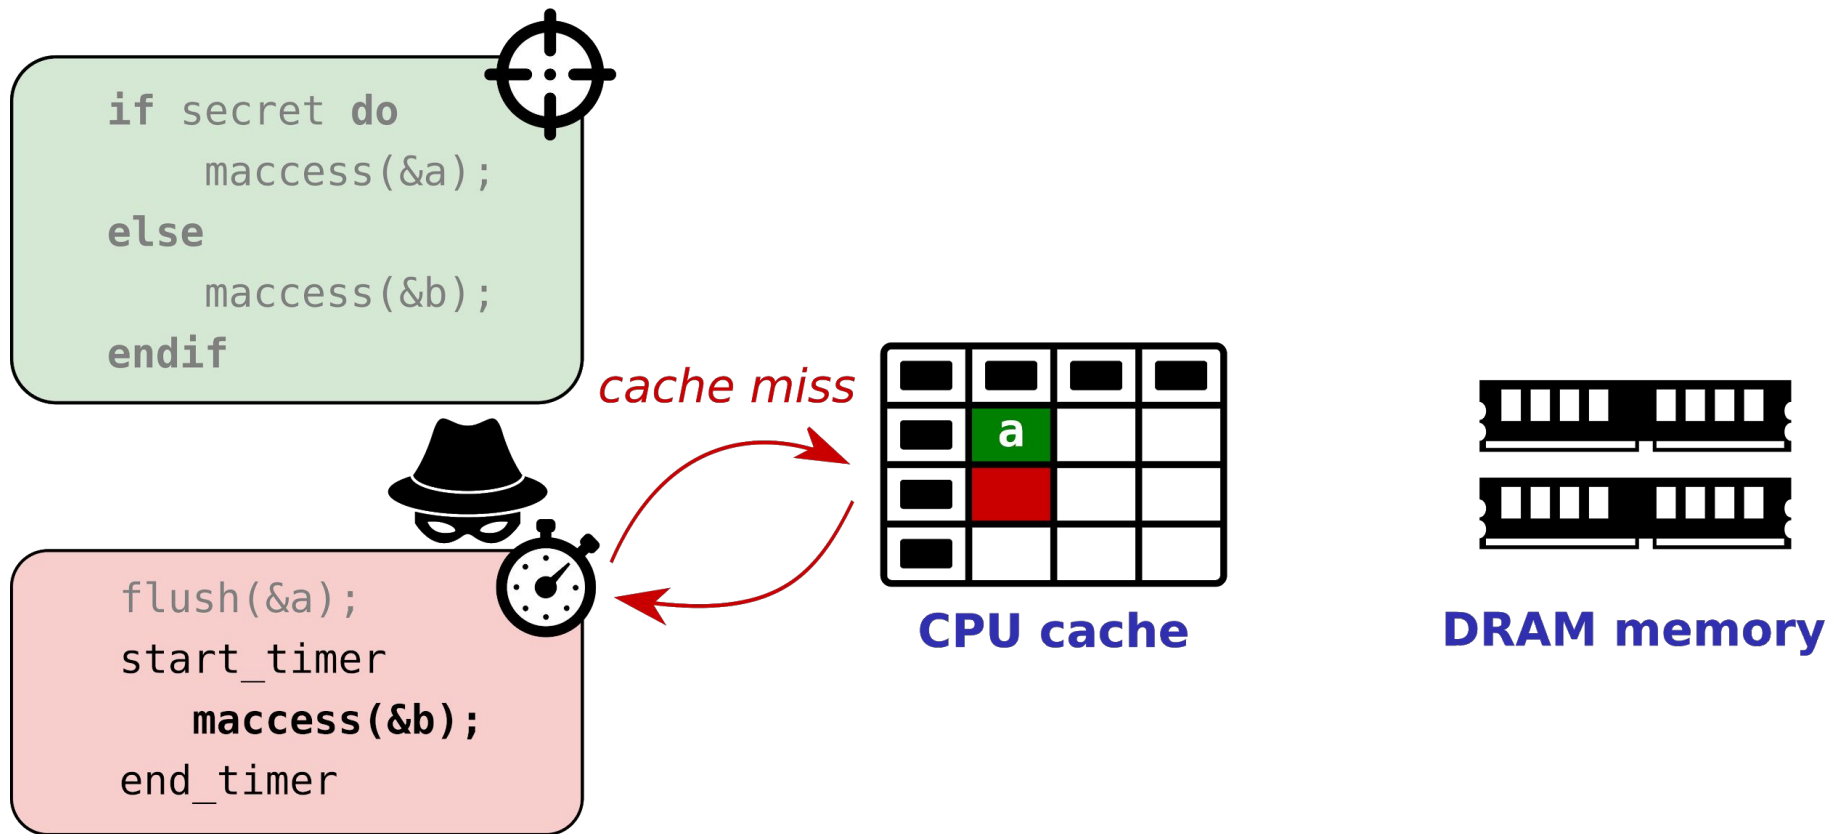
Flush+Reload: Cache timing attacks on shared memory



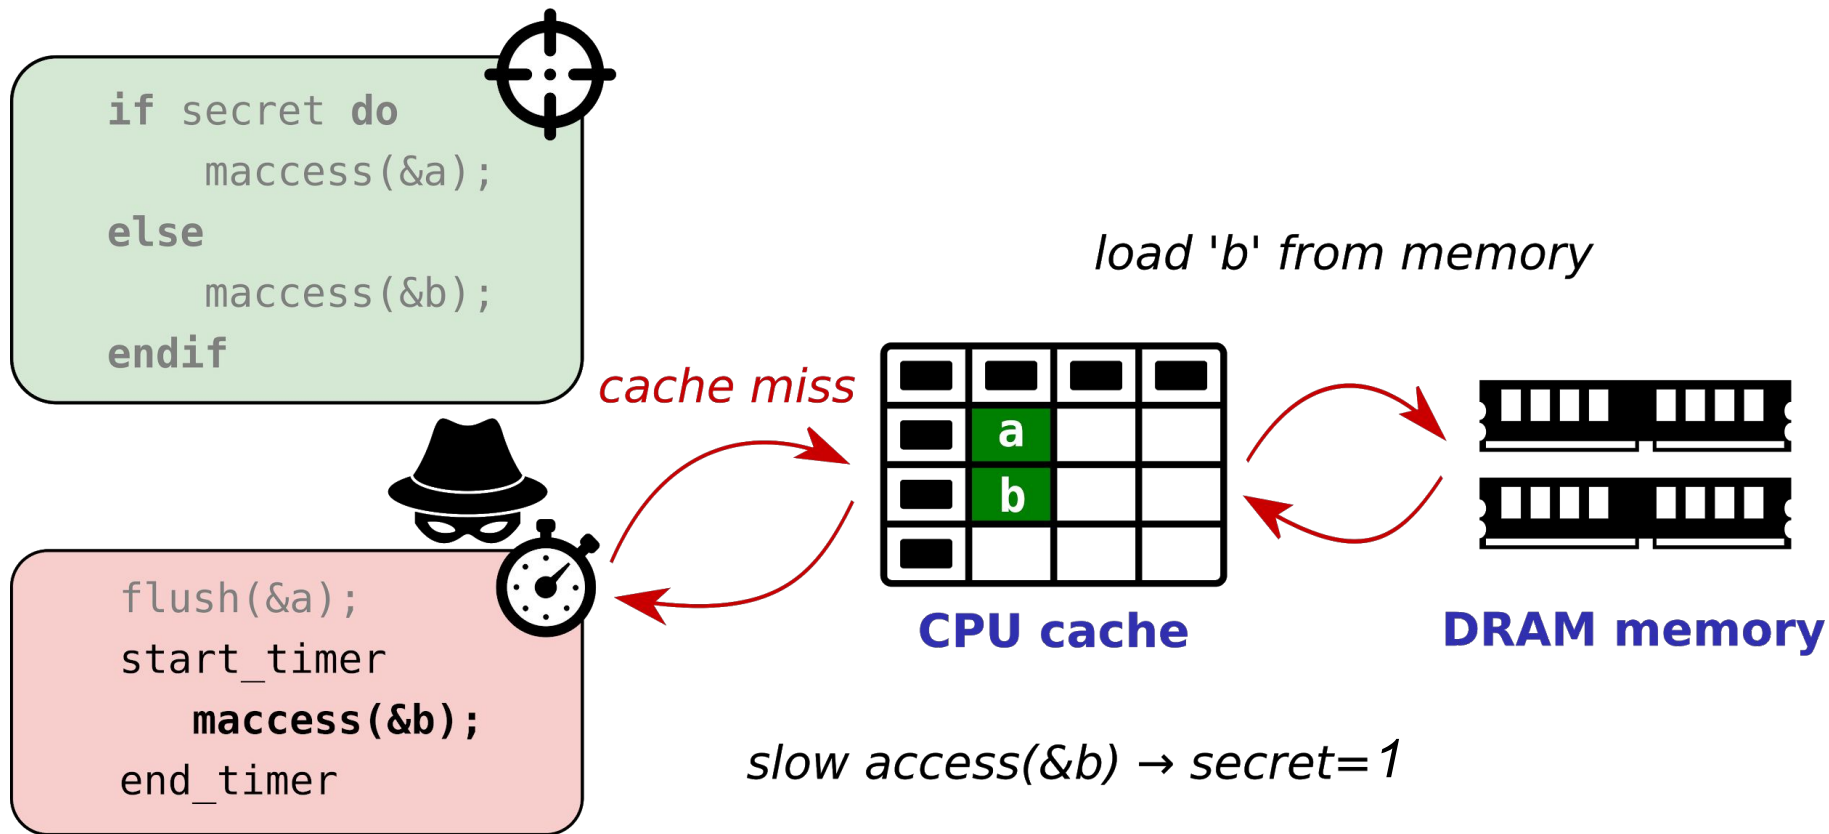
Flush+Reload: Cache timing attacks on shared memory



Flush+Reload: Cache timing attacks on shared memory



Flush+Reload: Cache timing attacks on shared memory



Flush+Reload Limitations

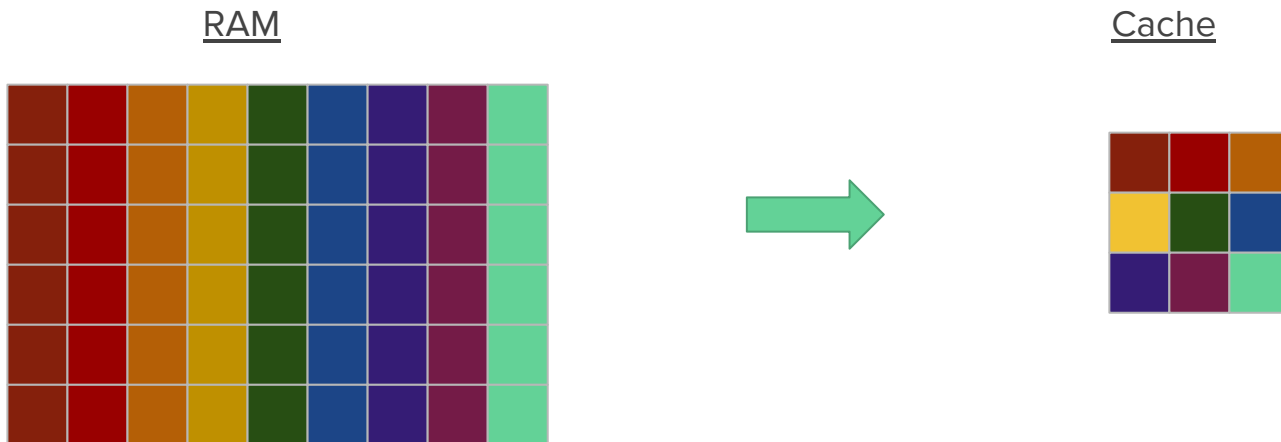
- Very **reliable and easy** attack
- But requires **shared memory** between victim and attacker apps
 - otherwise attacker cannot *flush* victim cache lines

→ **Generic** attack: does *not require knowledge* of internal cache organization, but only applicable in *limited scenarios* (when victim and attacker share memory)

Where in the cache should we store the contents of a memory location?

Direct mapping: concept

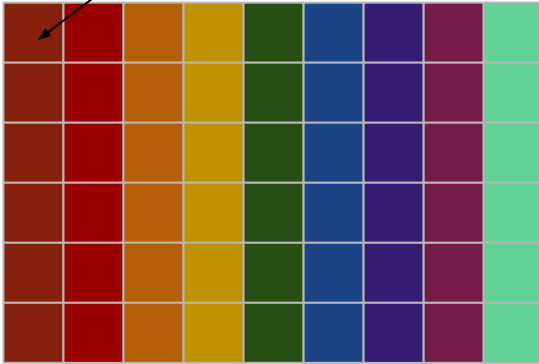
Notice: memory block index % 9 = cache block index



Direct mapping: concept

Each cache block has multiple addresses!

RAM



size of a block: 2^n bytes

=>

Address in same block share all but last n bits

Cache



size of a block: 2^n bytes \Rightarrow Address in same block share all but last n bits

Why?

Take $n = 3$

\Rightarrow 8 memory (byte)-addresses per block

Notice #1: shared bits = *memory block index* in binary!!

Notice #2: *last n bits* are in index within the block

Block index	Memory byte-address range	Binary range
0	[0, 7]	[00 000 - 00 111]
1	[8, 15]	[01 000 - 01 111]
2	[16, 23]	[10 000 - 10 111]
...

size of a block: 2^n bytes \Rightarrow Address in same block share all but last n bits

Why?

Take $n = 3$

\Rightarrow 8 memory (byte)-addresses per block

Remember
memory block index % cache
size (in blocks) = cache block
index

Notice #1: shared bits = *memory block index* in binary!!

Notice #2: *last n bits* are in index within the block

Block index	Memory byte-address range	Binary range
0	[0, 7]	[00 000 - 00 111]
1	[8, 15]	[01 000 - 01 111]
2	[16, 23]	[10 000 - 10 111]
...

binary(0xabc) =

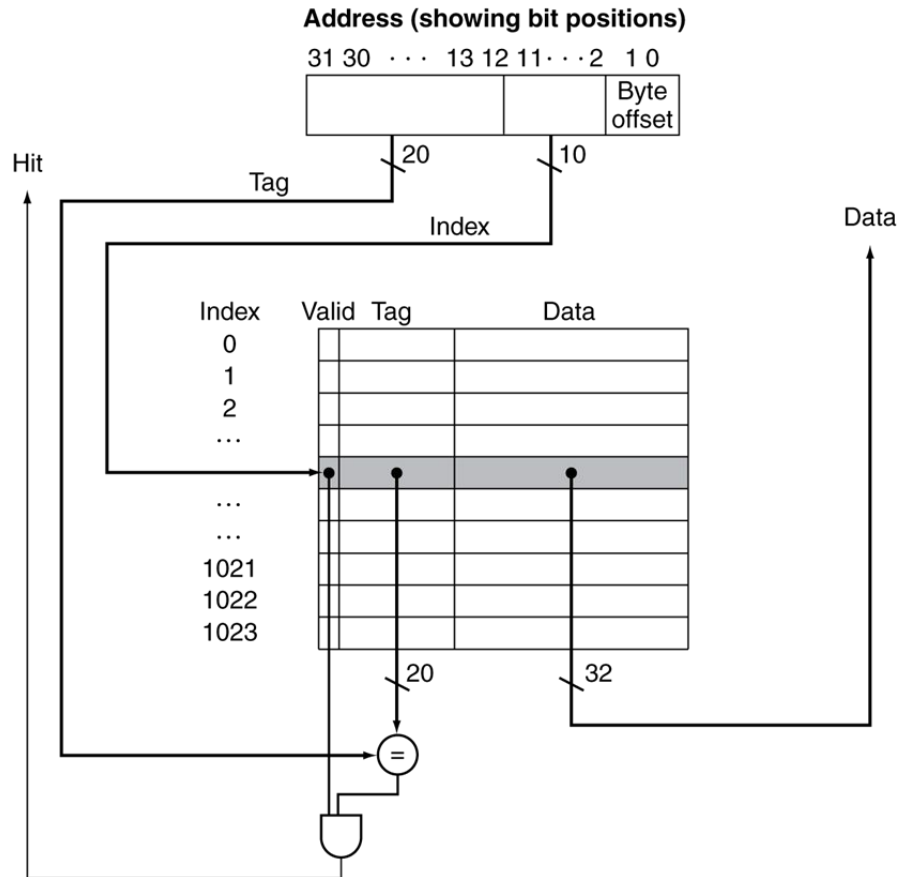
$\{1010\} \{10111\} \{100\}$
= 10 = 23 = 4

Assume:
Cache size: 32 blocks (2^5)
Block size: 8 bytes (2^3)

<u>Index</u>	Valid	Tag	Data
0	0	???	???
...	
10111 (23)	1	1010	8 bytes (0xab8 - 0xabf)
...	
11111 (31)	???	???	???

Notice:

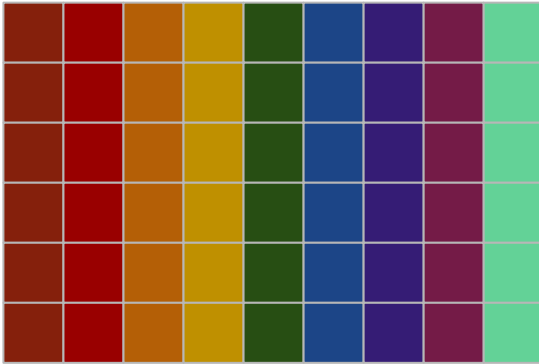
Index determines where to look
Tag determines hit or miss




Set-associativity: concept

Notice: each memory block can be located in n different **ways**

RAM



4-way Cache



Set-associativity: concept


Notice
 $n = 1$: direct mapping!

Notice: each memory block can be located in n different **ways**

RAM



4-way Cache




Notice:

Index determines where to look
BUT: multiple (#ways) possibilities

Check all:

Tag determines hit or miss



Fully associative cache

- n-way set associative cache
 - $n = \text{\#memory-blocks}$
- Best performance possible
 - only miss when a new unique address is accessed
 - any repeated request is a hit!

E.g., a cache with 8 blocks:

[illegible]

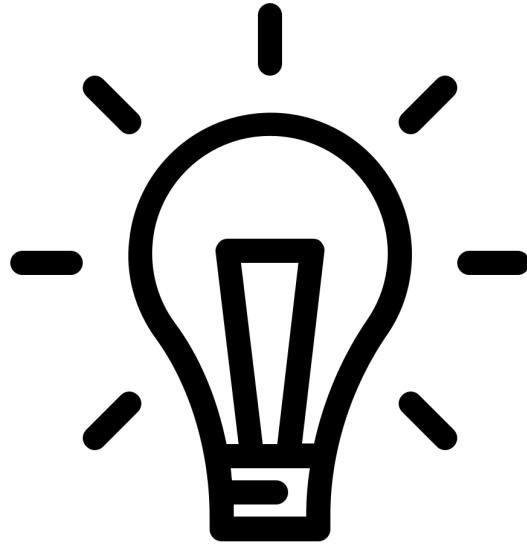
Fully associative cache

Basically, searching the entire cache without any indexing
– sequentially super slow!

Solution: search in different ways in parallel.
Need n comparators for any associative cache.


[illegible]






Can we leverage detailed knowledge of mapping schemes and cache collisions to mount more advanced cache timing attack?

Prime+Probe: Cache timing attacks across protection domains




```
if secret do
    maccess(&a);
else
    maccess(&b);
endif
```


'a' is **not** accessible
to attacker



CPU cache

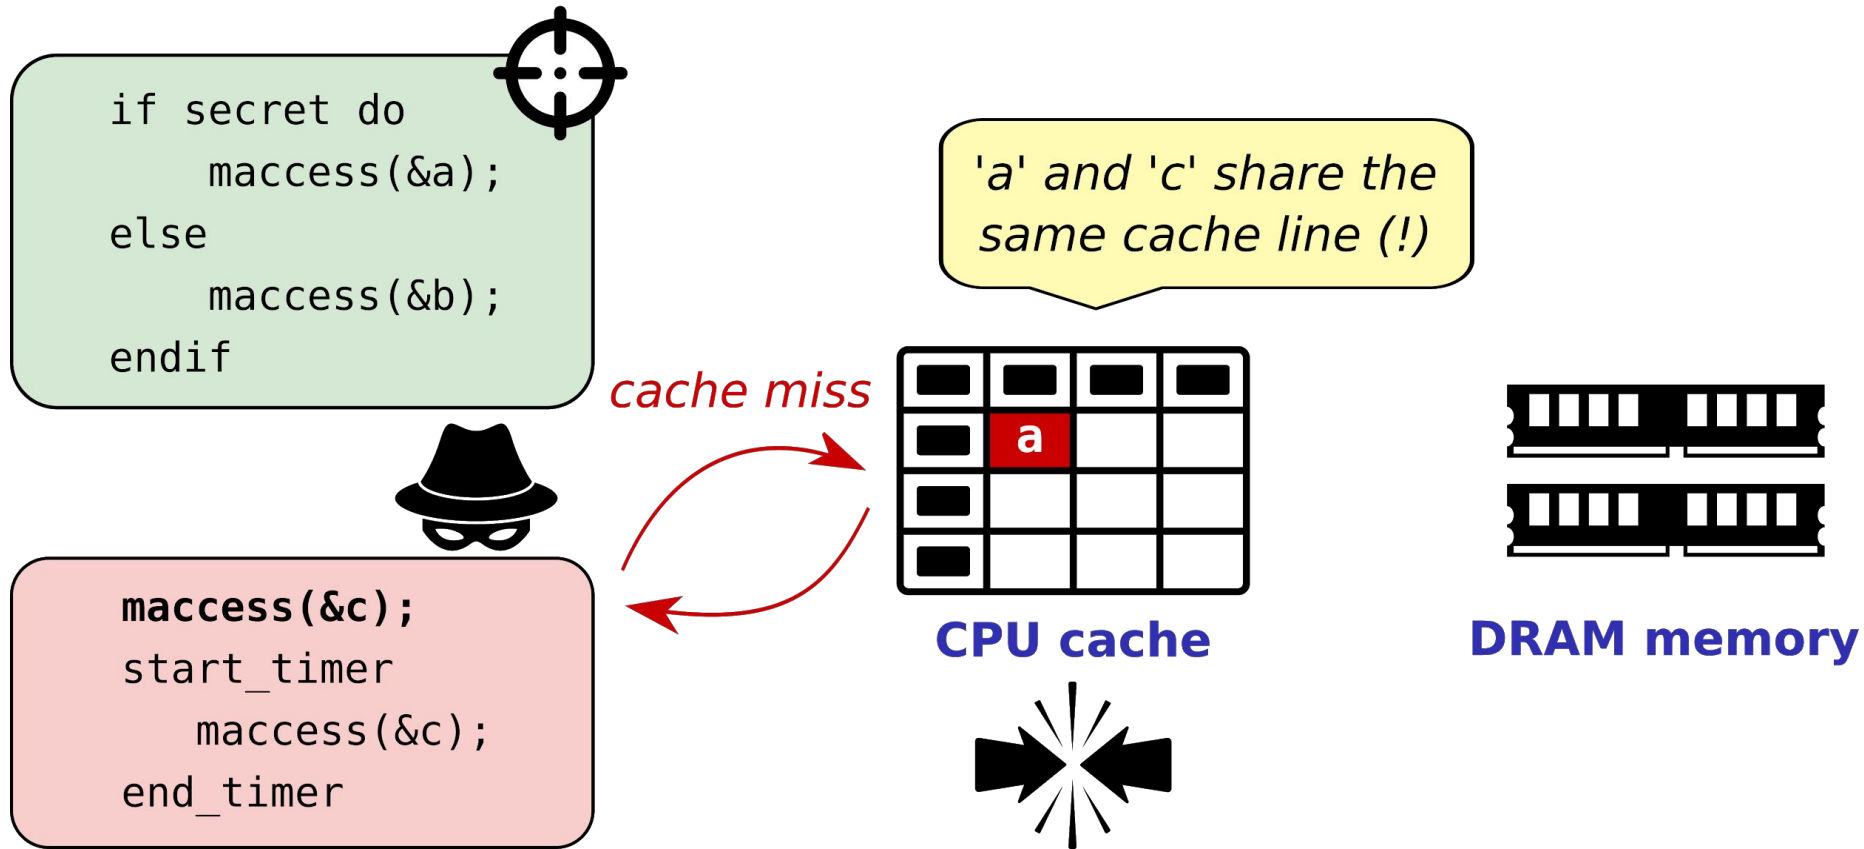


DRAM memory

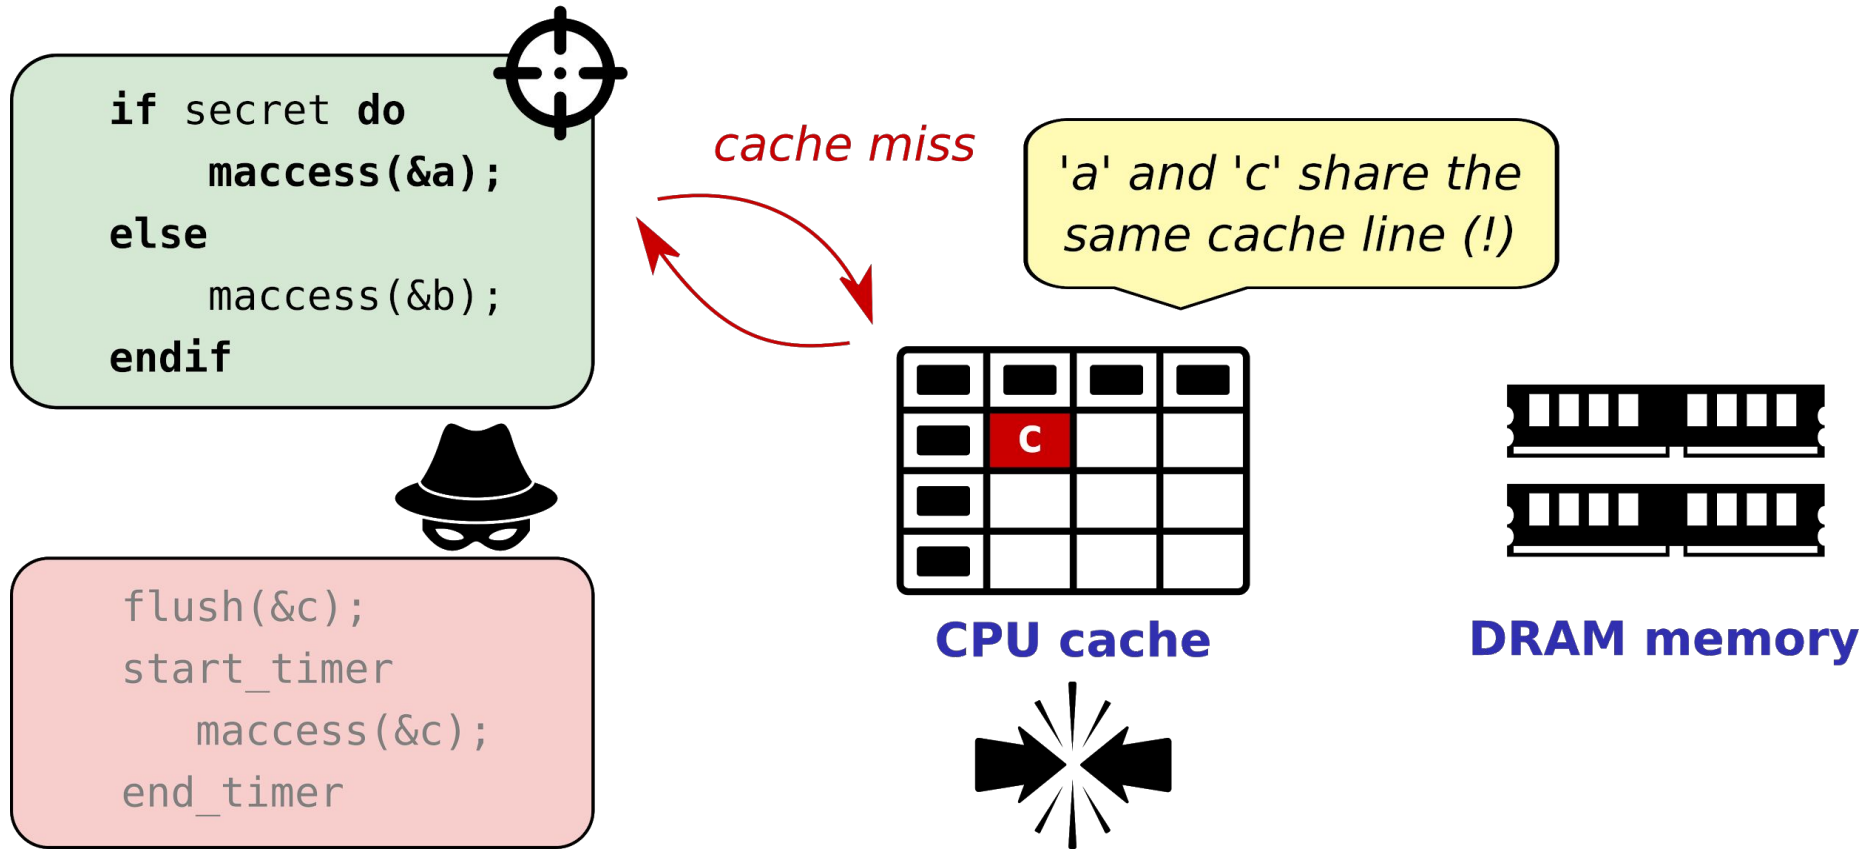


```
maccess(&c);
start_timer
    maccess(&c);
end_timer
```

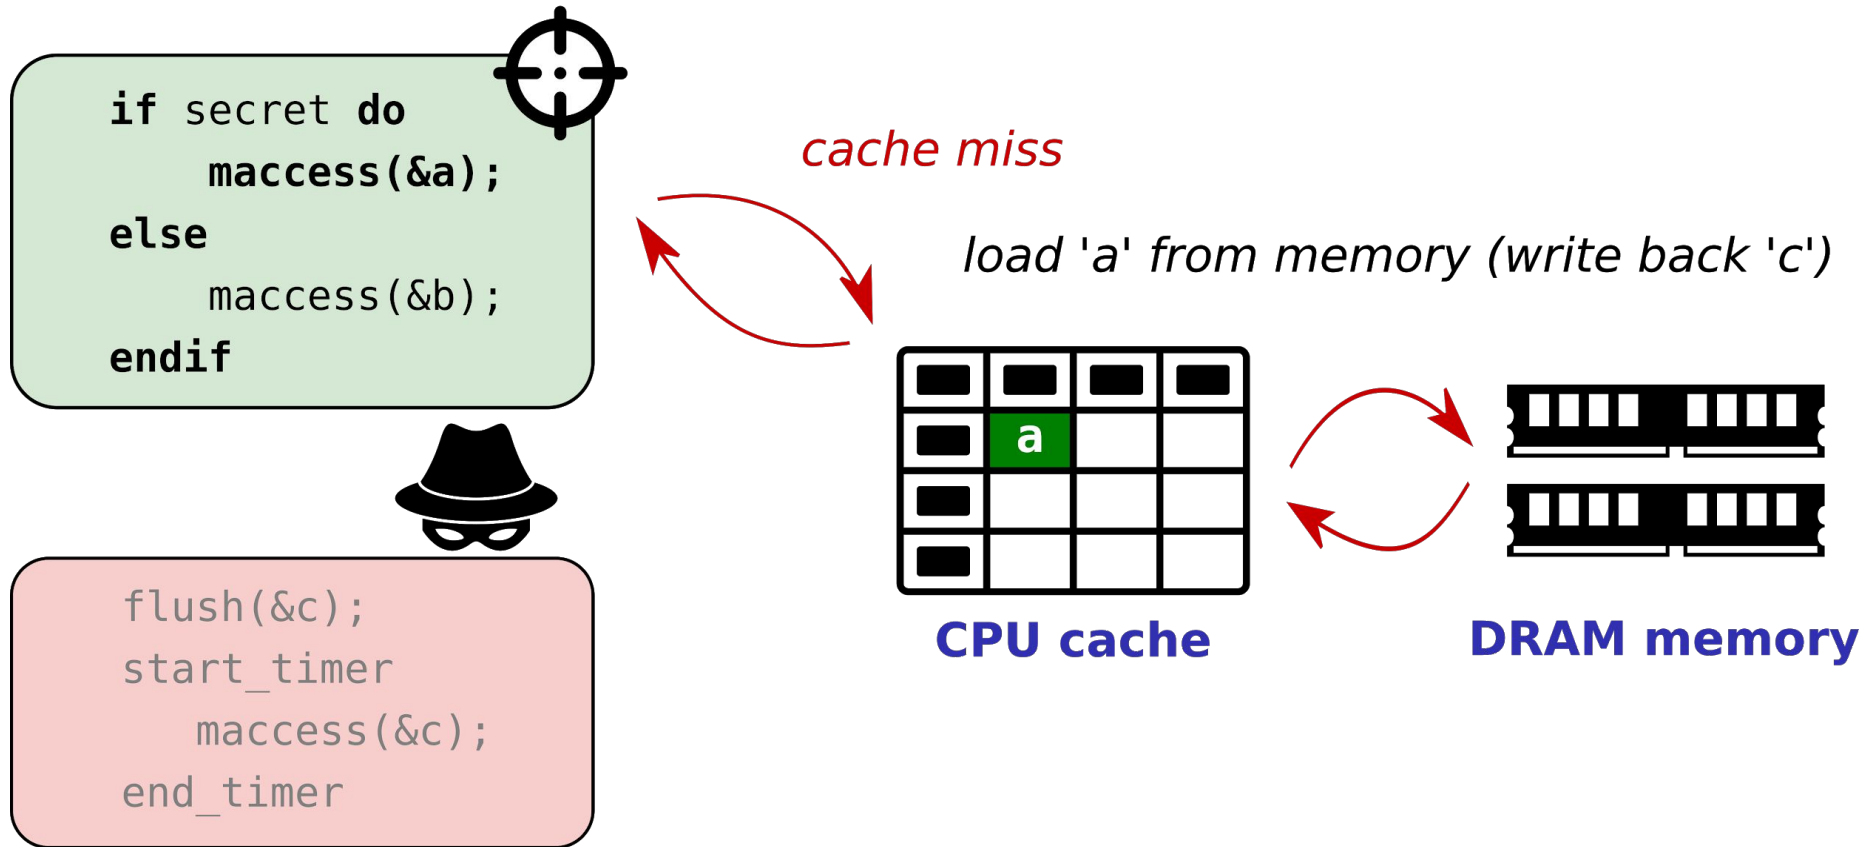

Prime+Probe: Cache timing attacks across protection domains



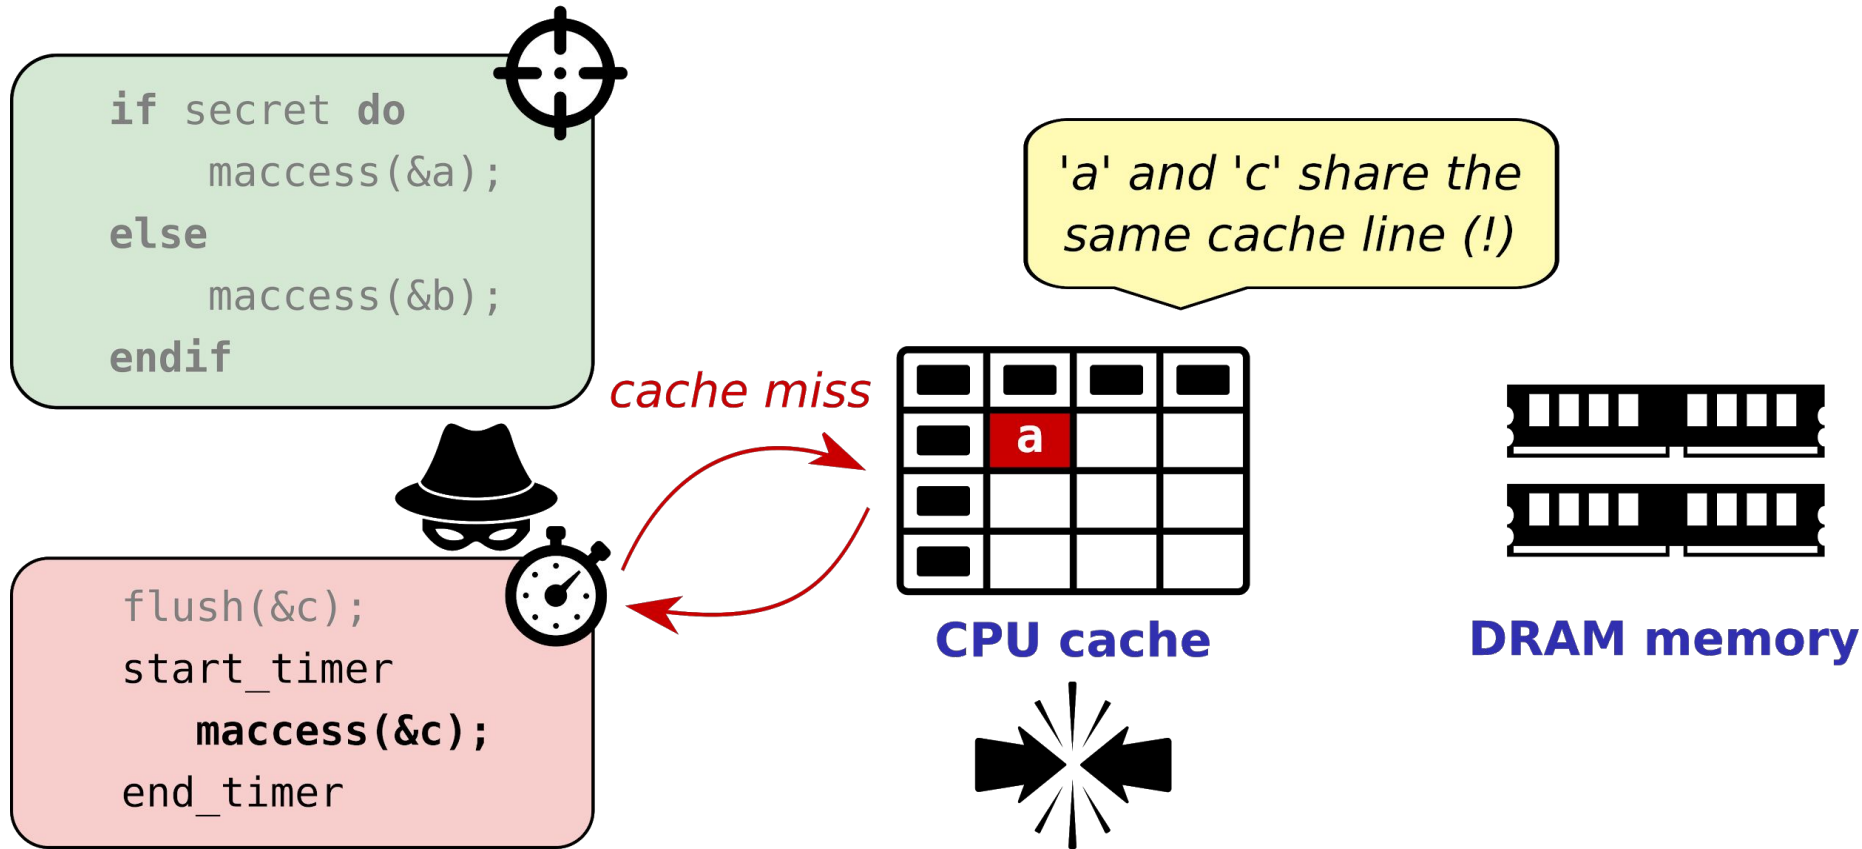
Prime+Probe: Cache timing attacks across protection domains



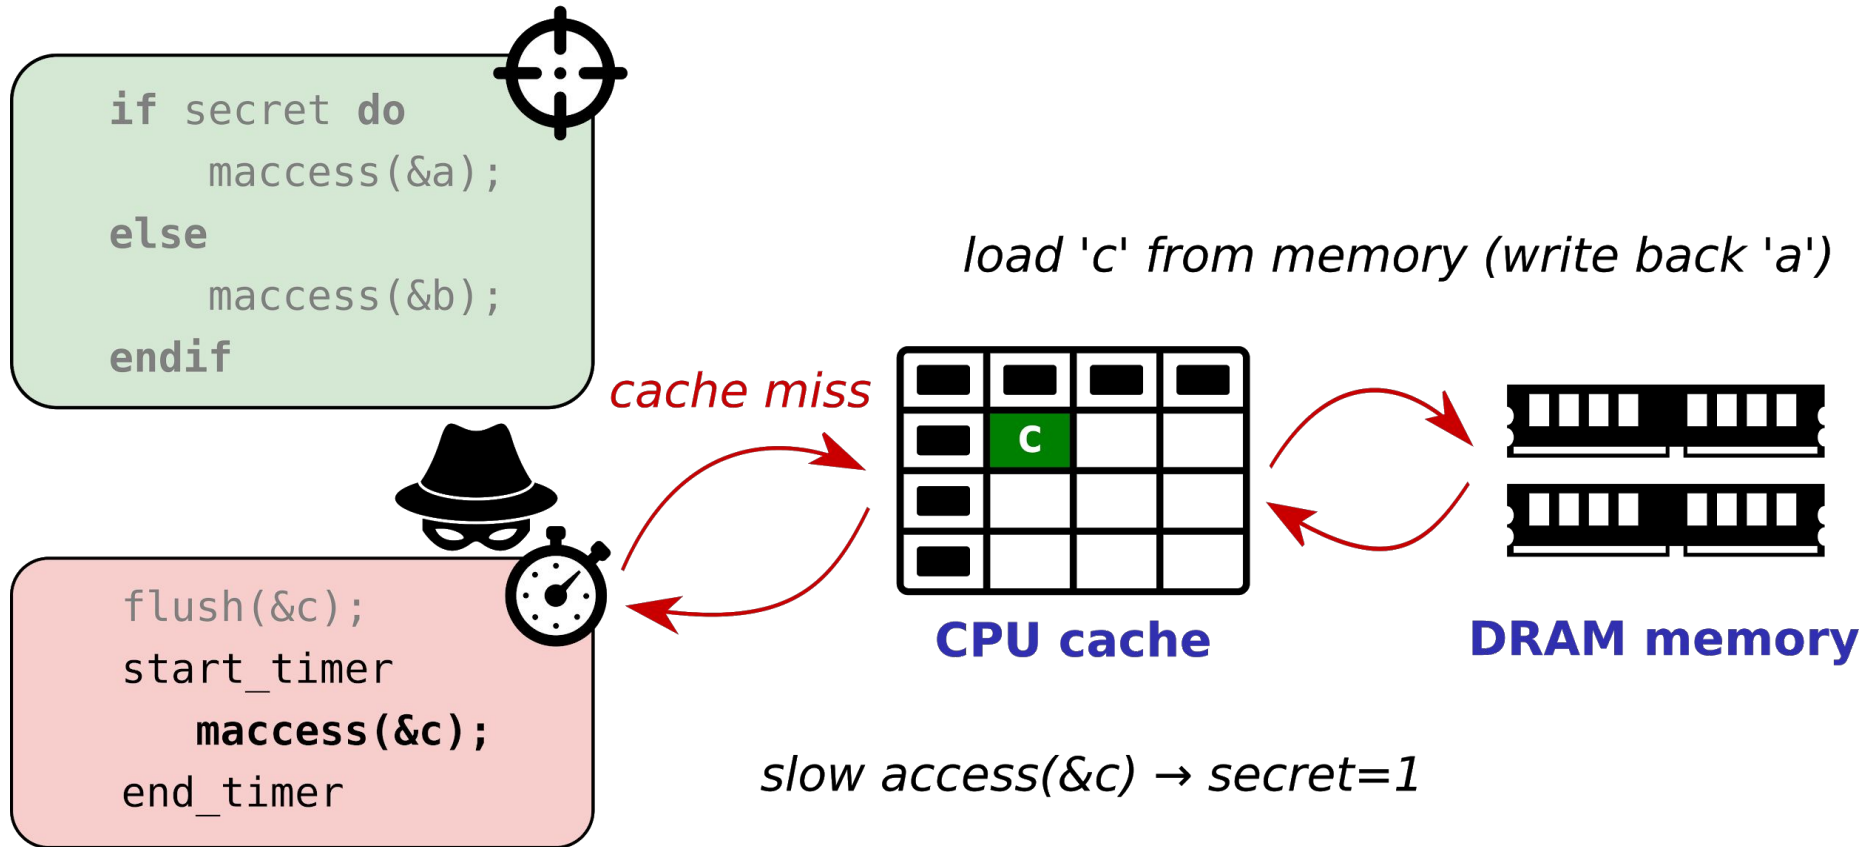
Prime+Probe: Cache timing attacks across protection domains



Prime+Probe: Cache timing attacks across protection domains



Prime+Probe: Cache timing attacks across protection domains



Appendix

Example 1: direct mapping 1-word blocks

Word address references:
22, 26, 22, 27, 3, 27, 53, 55

010110 → miss
011010 → miss
010110 → hit
011011 → miss
000011 → miss
011011 → miss
110101 → miss
110111 → miss

Index	V	Tag	Data
000	0		
001	0		
010	0		
011	0		
100	0		
101	0		
110	1	010	22
111	0		

Index	V	Tag	Data
000	0		
001	0		
010	1	011	26
011	0		
100	0		
101	0		
110	1	010	22
111	0		

Index	V	Tag	Data
000	0		
001	0		
010	1	011	26
011	1	011	27
100	0		
101	0		
110	1	010	22
111	0		

Index	V	Tag	Data
000	0		
001	0		
010	1	011	26
011	1	000	3
100	0		
101	0		
110	1	010	22
111	0		

Index	V	Tag	Data
000	0		
001	0		
010	1	011	26
011	1	011	27
100	0		
101	0		
110	1	010	22
111	0		

Index	V	Tag	Data
000	0		
001	0		
010	1	011	26
011	1	011	27
100	0		
101	0	110	53
110	1	010	22
111	0		

Index	V	Tag	Data
000	0		
001	0		
010	1	011	26
011	1	011	27
100	0		
101	0	110	53
110	1	010	22
111	1	110	55

Example 2: direct mapping 4-word blocks

Word address references:
22, 26, 22, 27, 3, 27, 53, 55

010110 → 101 → miss
011010 → 110 → miss
010110 → 101 → hit
011011 → 110 → hit
000011 → 000 → miss
011011 → 110 → hit
110101 → 110 → miss
110111 → 110 → hit

Index			Data			
			W1	W2	W3	W4
000	1					
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	0					
111	0					

Index			Data			
			W1	W2	W3	W4
000	0					
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	011	24	25	26	27
111	0					

Index			Data			
			W1	W2	W3	W4
000	1	000	0	1	2	3
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	011	24	25	26	27
111	0					

Index			Data			
			W1	W2	W3	W4
000	1	000	0	1	2	3
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	110	52	53	54	55
111	0					

Example 2: direct mapping 4-word blocks

Word address references:
22, 26, 22, 27, 3, 27, 53, 55

010110 → 101 → miss
011010 → 110 → miss
010110 → 101 → hit
011011 → 110 → hit
000011 → 000 → miss
011011 → 110 → hit
110101 → 110 → miss
110111 → 110 → hit

Miss rate
decreased.
Why?

Index			Data			
			W1	W2	W3	W4
000	1					
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	0					
111	0					

Index			Data			
			W1	W2	W3	W4
000	1	000	0	1	2	3
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	011	24	25	26	27
111	0					

Index			Data			
			W1	W2	W3	W4
000	0					
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	011	24	25	26	27
111	0					

Index			Data			
			W1	W2	W3	W4
000	1	000	0	1	2	3
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	110	52	53	54	55
111	0					

Example 2: direct mapping 4-word blocks

Word address references:
22, 26, 22, 27, 3, 27, 53, 55

010110 → 101 → miss
011010 → 110 → miss
010110 → 101 → hit
011011 → 110 → hit
000011 → 000 → miss
011011 → 110 → hit
110101 → 110 → miss
110111 → 110 → hit

**Miss rate
decreased.
Why?**

**But miss penalty
increased.
Why?**

Index			Data			
			W1	W2	W3	W4
000	1					
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	0					
111	0					

Index			Data			
			W1	W2	W3	W4
000	1	000	0	1	2	3
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	011	24	25	26	27
111	0					

Index			Data			
			W1	W2	W3	W4
000	0					
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	011	24	25	26	27
111	0					

Index			Data			
			W1	W2	W3	W4
000	1	000	0	1	2	3
001	0					
010	0					
011	0					
100	0					
101	1	010	20	21	22	23
110	1	110	52	53	54	55
111	0					

Example 3: 2-way set-associative 1-word blocks

Word address references:
22, 26, 22, 27, 3, 27, 53, 55

010110 → miss
011010 → miss
010110 → hit
011011 → miss
000011 → miss
011011 → hit
110101 → miss
110111 → miss

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22		
11	0				

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	1	0110	27		


Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	1	1101	53		
10	1	0101	22	0110	26
11	1	0110	27	0000	3

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	0				

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	1	0110	27	0000	3

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	1	1101	53		
10	1	0101	22	0110	26
11	1	0110	27	1101	55

Replacement rule:
Least Recently Used



Example 3: 2-way set-associative 1-word blocks

Word address references:
22, 26, 22, 27, 3, 27, 53, 55

010110 → miss
011010 → miss
010110 → hit
011011 → miss
000011 → miss
011011 → hit
110101 → miss
110111 → miss

Miss rate
decreased.
Why?

Trade-off?

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22		
11	0				

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	1	0110	27		

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	1	1101	53		
10	1	0101	22	0110	26
11	1	0110	27	0000	3

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	0				

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	1	0110	27	0000	3

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	1	1101	53		
10	1	0101	22	0110	26
11	1	0110	27	1101	55

Example 3: 2-way set-associative 1-word blocks

Word address references:
22, 26, 22, 27, 3, 27, 53, 55

010110 → miss
011010 → miss
010110 → hit
011011 → miss
000011 → miss
011011 → hit
110101 → miss
110111 → miss

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22		
11	0				

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	1	0110	27		

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	1	1101	53		
10	1	0101	22	0110	26
11	1	0110	27	0000	3

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	0				

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	0				
10	1	0101	22	0110	26
11	1	0110	27	0000	3

Set	V	Tag0	Data0	Tag1	Data1
00	0				
01	1	1101	53		
10	1	0101	22	0110	26
11	1	0110	27	1101	55

Miss rate
decreased.
Why?

Hit time
increased!
We search longer

